

```
## Laptop Price Prediction for SmartTech Co
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

## Data Exploration and Understanding:

```
df=pd.read_csv(r"D:\ml\laptop.csv") # reading dataset
```

```
df.sample(5) # displays 5 random rows
```

|      | Unnamed: 0.1 | Unnamed: 0 | Company | TypeName  | Inches | \ |
|------|--------------|------------|---------|-----------|--------|---|
| 763  | 763          | 763.0      | Asus    | Ultrabook | 13.3   |   |
| 466  | 466          | 466.0      | Acer    | Notebook  | 15.6   |   |
| 705  | 705          | 705.0      | Dell    | Notebook  | 15.6   |   |
| 1240 | 1240         | 1240.0     | Lenovo  | Notebook  | 15.6   |   |
| 634  | 634          | 634.0      | Asus    | Notebook  | 15.6   |   |

|        |           |          | ScreenResolution |                               |
|--------|-----------|----------|------------------|-------------------------------|
| Cpu    | Ram       | \        |                  |                               |
| 763    | IPS Panel | Quad HD+ | 3200x1800        | Intel Core i5 7200U           |
| 2.5GHz | 8GB       |          |                  |                               |
| 466    |           |          | 1366x768         | Intel Core i3 6006U           |
| 2GHz   | 4GB       |          |                  |                               |
| 705    |           | Full HD  | 1920x1080        | Intel Core i5 7200U           |
| 2.5GHz | 8GB       |          |                  |                               |
| 1240   |           | Full HD  | 1920x1080        | AMD A12-Series 9720P          |
| 3.6GHz | 6GB       |          |                  |                               |
| 634    |           |          | 1366x768         | Intel Celeron Dual Core N3350 |
| 1.1GHz | 8GB       |          |                  |                               |

|            | Memory    |                          | Gpu        | OpSys  | Weight |
|------------|-----------|--------------------------|------------|--------|--------|
| Price      |           |                          |            |        |        |
| 763        | 256GB SSD | Intel HD Graphics 620    | Windows 10 | 1.2kg  |        |
| 60153.1200 |           |                          |            |        |        |
| 466        | 500GB HDD | Nvidia GeForce GTX 940MX | Windows 10 | 2.2kg  |        |
| 24988.3200 |           |                          |            |        |        |
| 705        | 256GB SSD | Intel HD Graphics 620    | Windows 10 | 2.18kg |        |
| 42357.6000 |           |                          |            |        |        |
| 1240       | 256GB SSD | AMD Radeon 530           | Windows 10 | 2.2kg  |        |
| 31838.5296 |           |                          |            |        |        |
| 634        | 1TB HDD   | Intel HD Graphics 500    | Windows 10 | 2kg    |        |
| 21258.7200 |           |                          |            |        |        |

```
df.columns
```

```
Index(['Unnamed: 0.1', 'Unnamed: 0', 'Company', 'TypeName', 'Inches',
      'ScreenResolution', 'Cpu', 'Ram', 'Memory', 'Gpu', 'OpSys',
      'Weight',
      'Price'],
      dtype='object')
```

*# Removing Unwanted columns*

```
df.drop(columns=["Unnamed: 0.1", 'Unnamed: 0'], inplace=True)
```

```
df.columns
```

```
Index(['Company', 'TypeName', 'Inches', 'ScreenResolution', 'Cpu',
      'Ram',
      'Memory', 'Gpu', 'OpSys', 'Weight', 'Price'],
      dtype='object')
```

```
print(f"{df.shape}") # There are 1303 rows and 11 columns
```

```
(1303, 11)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1303 entries, 0 to 1302
```

```
Data columns (total 11 columns):
```

| #  | Column           | Non-Null Count | Dtype   |
|----|------------------|----------------|---------|
| 0  | Company          | 1273 non-null  | object  |
| 1  | TypeName         | 1273 non-null  | object  |
| 2  | Inches           | 1273 non-null  | object  |
| 3  | ScreenResolution | 1273 non-null  | object  |
| 4  | Cpu              | 1273 non-null  | object  |
| 5  | Ram              | 1273 non-null  | object  |
| 6  | Memory           | 1273 non-null  | object  |
| 7  | Gpu              | 1273 non-null  | object  |
| 8  | OpSys            | 1273 non-null  | object  |
| 9  | Weight           | 1273 non-null  | object  |
| 10 | Price            | 1273 non-null  | float64 |

```
dtypes: float64(1), object(10)
```

```
memory usage: 112.1+ KB
```

```
df.isnull().sum() # there are 30 null values in each row
```

|                  |    |
|------------------|----|
| Company          | 30 |
| TypeName         | 30 |
| Inches           | 30 |
| ScreenResolution | 30 |
| Cpu              | 30 |
| Ram              | 30 |
| Memory           | 30 |
| Gpu              | 30 |

```
OpSys      30
Weight     30
Price      30
dtype: int64
```

*# As there are equal null values in each column Dropping the rows rows which contains null values*

```
df.dropna(inplace=True)
```

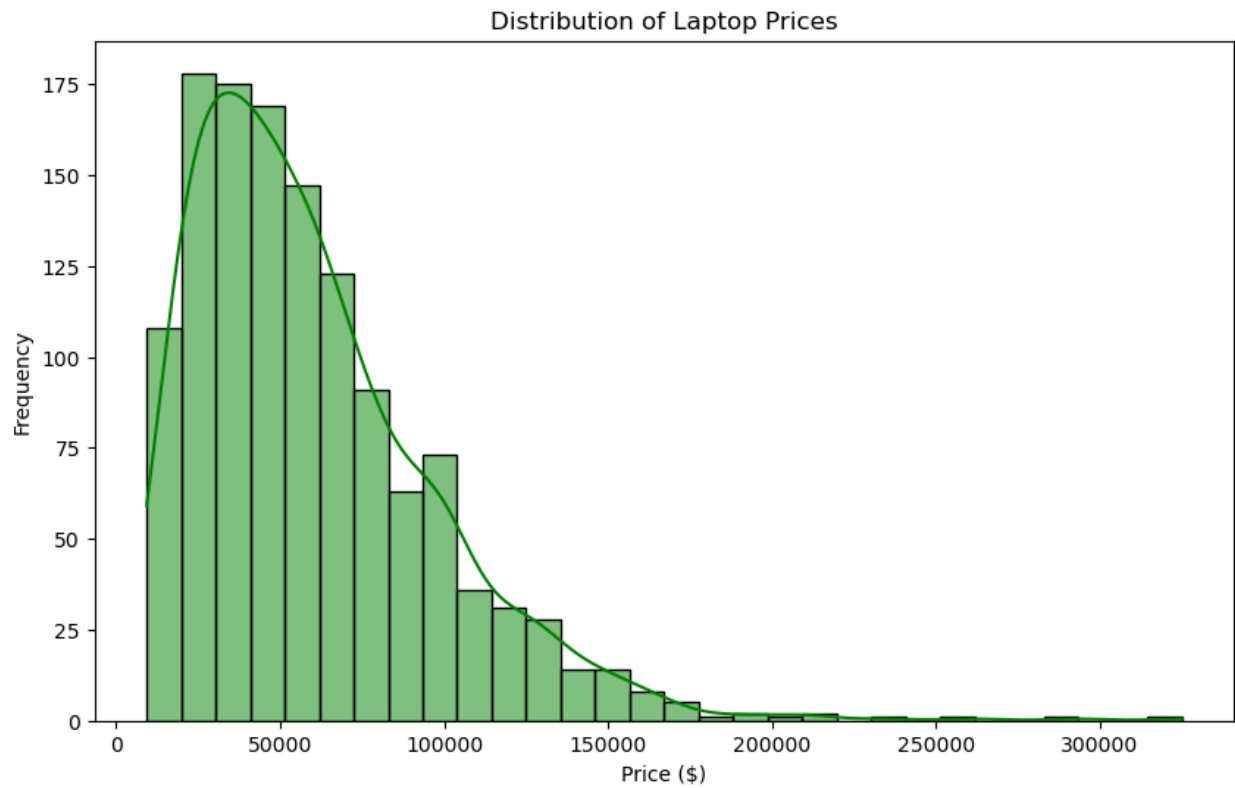
```
df.isnull().sum() # There are no null values
```

```
Company      0
TypeName     0
Inches       0
ScreenResolution  0
Cpu          0
Ram          0
Memory       0
Gpu          0
OpSys        0
Weight       0
Price        0
dtype: int64
```

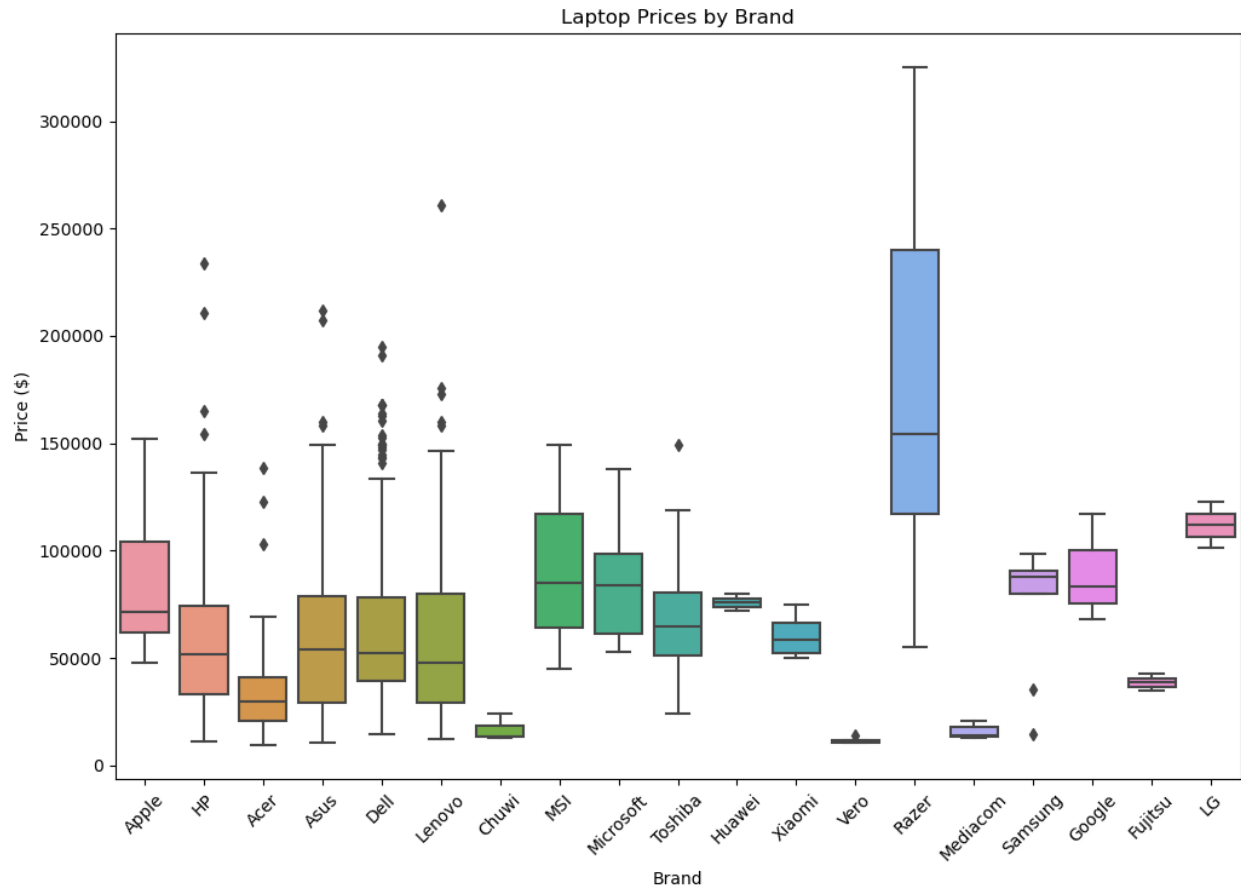
## Data Visualization

*# Visualize trends in laptop prices*

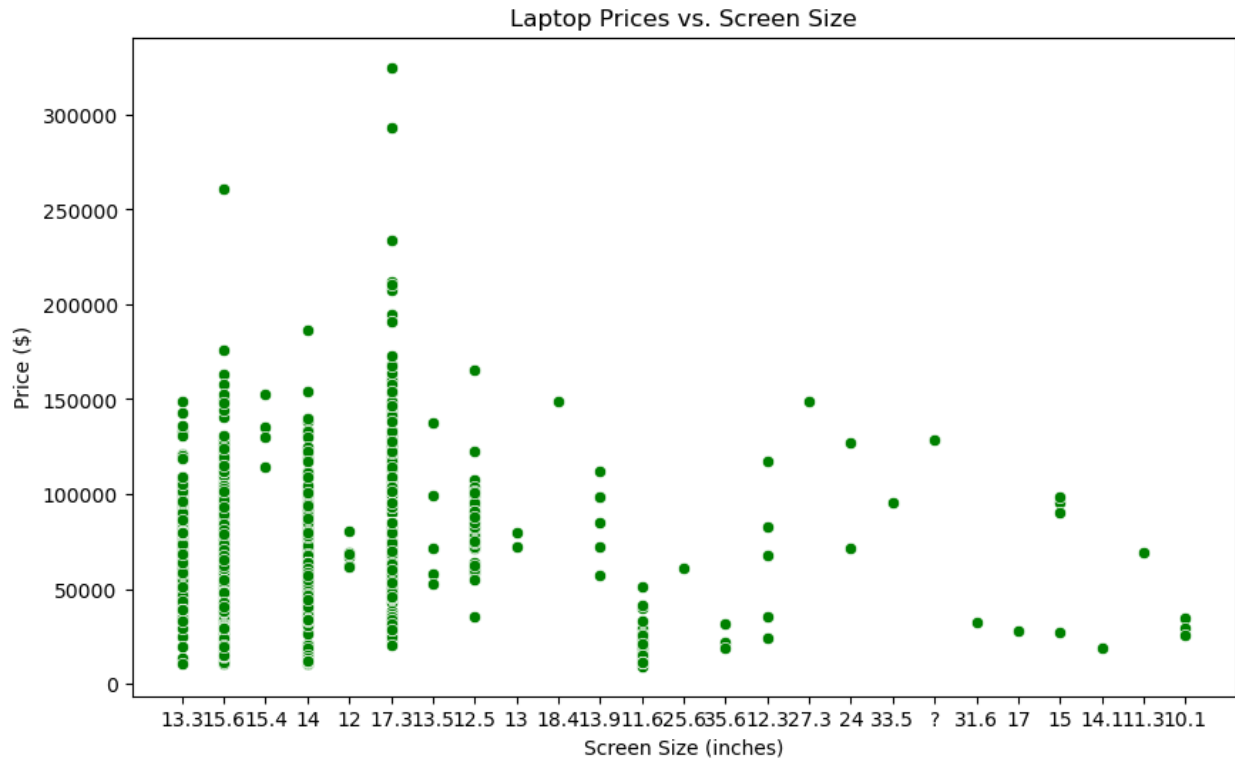
```
plt.figure(figsize=(10, 6))
sns.histplot(df['Price'], bins=30, kde=True, color='green')
plt.title('Distribution of Laptop Prices')
plt.xlabel('Price ($)')
plt.ylabel('Frequency')
plt.show()
```



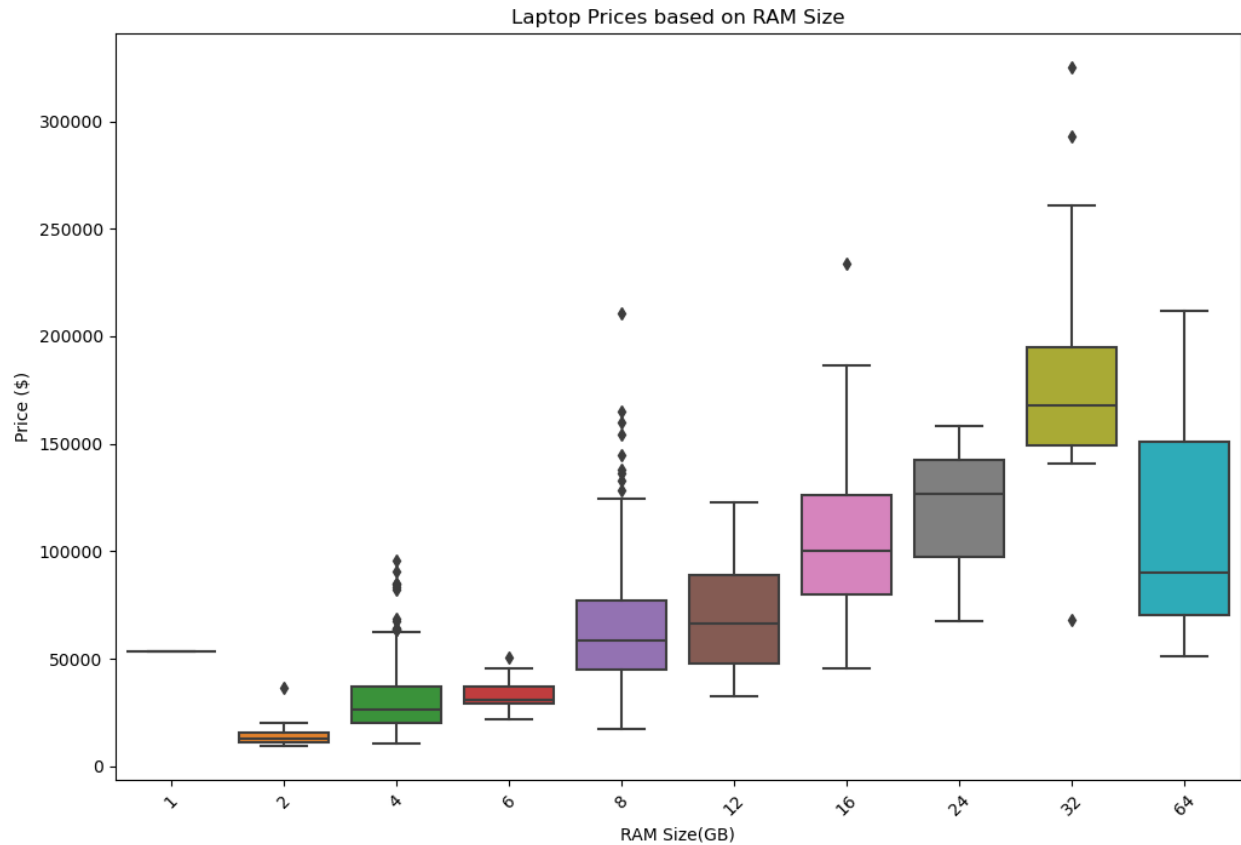
```
# Boxplot of laptop prices by brand
plt.figure(figsize=(12, 8))
sns.boxplot(x='Company', y='Price', data=df)
plt.title('Laptop Prices by Brand')
plt.xlabel('Brand')
plt.ylabel('Price ($)')
plt.xticks(rotation=45)
plt.show()
```




```
# Scatter plot of laptop prices against screen size
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Inches', y='Price', data=df, color='green')
plt.title('Laptop Prices vs. Screen Size')
plt.xlabel('Screen Size (inches)')
plt.ylabel('Price ($)')
plt.show()
```



```
# Clean RAM Column
# Visualization of RAM
df['Ram'] = df['Ram'].str.replace('GB', '').astype(int)
plt.figure(figsize=(12, 8))
sns.boxplot(x='Ram', y='Price', data=df)
plt.title('Laptop Prices based on RAM Size')
plt.xlabel('RAM Size(GB)')
plt.ylabel('Price ($)')
plt.xticks(rotation=45)
plt.show()
```



```
# Visualize correlation matrix - Have to select only Numerical columns  
for Correlation  
numeric = df.select_dtypes(include=[float,int]).columns  
plt.figure(figsize=(12, 10))  
sns.heatmap(df[numeric].corr(), annot=True, cmap='coolwarm',  
fmt=".2f")  
plt.title('Correlation Matrix')  
plt.show()
```



## Data Preprocessing:

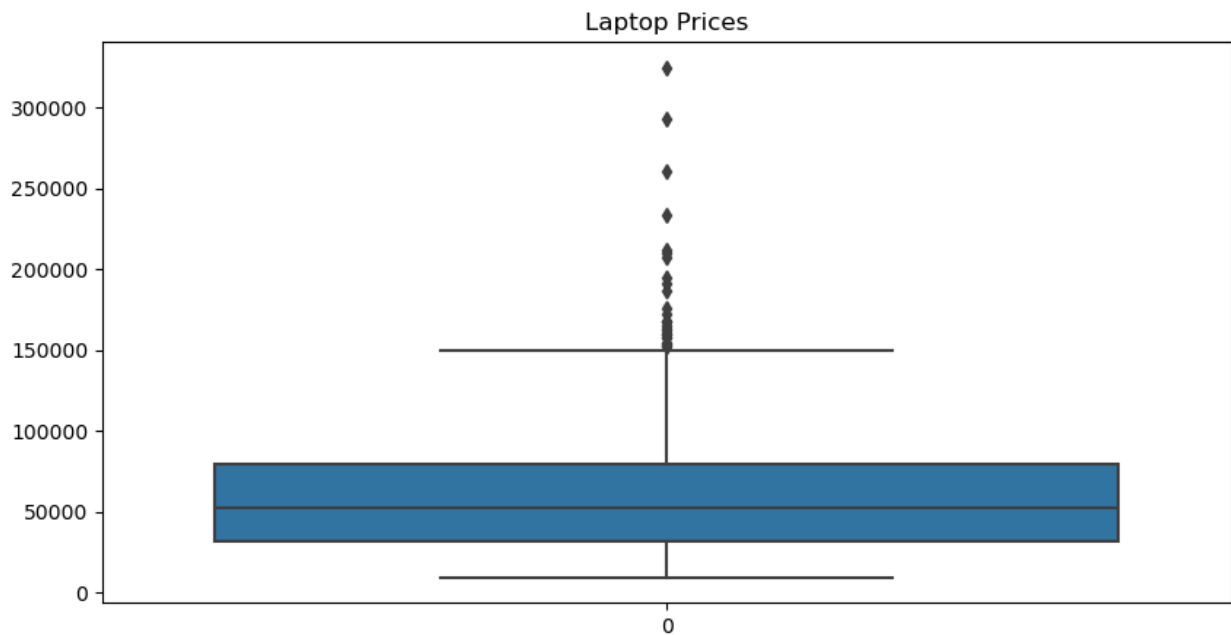
```
# Missing Values  
print(f"Missing Values:\n {df.isnull().sum()}")
```

```
Missing Values:  
Company      0  
TypeName     0  
Inches       0  
ScreenResolution  0  
Cpu          0  
Ram          0  
Memory       0  
Gpu          0  
OpSys        0  
Weight       0
```



```
Price          0
dtype: int64
```

```
# Identify Outlier in Price column using Box Plot
plt.figure(figsize=(10, 5))
sns.boxplot(df['Price'])
plt.title('Laptop Prices')
plt.show()
```



```
# Handling Outliers with IQR Method
def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column] >= lower_bound) & (df[column] <=
upper_bound)]

df = remove_outliers(df, 'Price')

print(f"Shape of dataset after handling Outliers {df.shape}")
Shape of dataset after handling Outliers (1245, 11)
```

Encode categorical variables

```
categorical = df.select_dtypes(include=[object]).columns
print(f"Categorical Columns {categorical}")
```

```

Categorical Columns Index(['Company', 'TypeName', 'Inches',
    'ScreenResolution', 'Cpu', 'Memory',
    'Gpu', 'OpSys', 'Weight'],
    dtype='object')

# Performing One-hot encoding on Categorical Variables
encoding = pd.get_dummies(df,columns=categorical,drop_first=True)

encoding.head()

```

|              | Ram | Price       | Company_Apple | Company_Asus | Company_Chui |  |
|--------------|-----|-------------|---------------|--------------|--------------|--|
| Company_Dell | \   |             |               |              |              |  |
| 0            | 8   | 71378.6832  | True          | False        | False        |  |
| False        |     |             |               |              |              |  |
| 1            | 8   | 47895.5232  | True          | False        | False        |  |
| False        |     |             |               |              |              |  |
| 2            | 8   | 30636.0000  | False         | False        | False        |  |
| False        |     |             |               |              |              |  |
| 3            | 16  | 135195.3360 | True          | False        | False        |  |
| False        |     |             |               |              |              |  |
| 4            | 8   | 96095.8080  | True          | False        | False        |  |
| False        |     |             |               |              |              |  |

|   | Company_Fujitsu | Company_Google | Company_HP | Company_Huawei | ... | \   |
|---|-----------------|----------------|------------|----------------|-----|-----|
| 0 | False           | False          | False      | False          | ... | ... |
| 1 | False           | False          | False      | False          | ... | ... |
| 2 | False           | False          | True       | False          | ... | ... |
| 3 | False           | False          | False      | False          | ... | ... |
| 4 | False           | False          | False      | False          | ... | ... |

|   | Weight_4.5kg | Weight_4.6kg | Weight_4kg | Weight_5.4kg | Weight_5.8kg |  |
|---|--------------|--------------|------------|--------------|--------------|--|
| \ |              |              |            |              |              |  |
| 0 | False        | False        | False      | False        | False        |  |
| 1 | False        | False        | False      | False        | False        |  |
| 2 | False        | False        | False      | False        | False        |  |
| 3 | False        | False        | False      | False        | False        |  |
| 4 | False        | False        | False      | False        | False        |  |

|   | Weight_6.2kg | Weight_7.2kg | Weight_8.23kg | Weight_8.4kg | Weight_? |  |
|---|--------------|--------------|---------------|--------------|----------|--|
| 0 | False        | False        | False         | False        | False    |  |
| 1 | False        | False        | False         | False        | False    |  |
| 2 | False        | False        | False         | False        | False    |  |
| 3 | False        | False        | False         | False        | False    |  |
| 4 | False        | False        | False         | False        | False    |  |

```
[5 rows x 531 columns]
```

Ensure the dataset is ready for model training

```
# Separating the features(X) and target(Y) Variables
X = encoding.drop(columns=["Price"])
y = encoding["Price"]

X.shape

(1245, 530)

y.shape

(1245,)
```

## Feature Engineering:

Extract meaningful features to enhance model performance

```
# Extracting CPU brand
df['CpuBrand'] = df['Cpu'].str.split().str[0]
df['CpuBrand'].unique()

array(['Intel', 'AMD', 'Samsung'], dtype=object)

# Extracting screen height and width from the screenresolution column
df['ScreenWidth'] = df['ScreenResolution'].str.extract(r'(\d+)x')[0].astype(int)
df['ScreenHeight'] = df['ScreenResolution'].str.extract(r'(\d+)x')[1].astype(int)

df[['ScreenResolution', 'ScreenWidth', 'ScreenHeight']].head()
```

|   |           |                |         | ScreenResolution | ScreenWidth | ScreenHeight |
|---|-----------|----------------|---------|------------------|-------------|--------------|
| 0 | IPS Panel | Retina Display |         | 2560x1600        | 2560        | 2560         |
| 1 |           |                |         | 1440x900         | 1440        | 1440         |
| 2 |           |                | Full HD | 1920x1080        | 1920        | 1920         |
| 3 | IPS Panel | Retina Display |         | 2880x1800        | 2880        | 2880         |
| 4 | IPS Panel | Retina Display |         | 2560x1600        | 2560        | 2560         |

Creation of New features

```
# Inches column should be Numeric and Have to handle conversion issues
df['Inches'] = pd.to_numeric(df['Inches'].str.replace('"', ''),
errors='coerce')
df.dropna(subset=['Inches'], inplace=True)
```

Encode Categorical Features

```
categorical_columns = df.select_dtypes(include=['object']).columns #
categorical_columns
```

```
Index(['Company', 'TypeName', 'ScreenResolution', 'Cpu', 'Memory',
      'Gpu',
      'OpSys', 'Weight', 'CpuBrand'],
      dtype='object')
```

```
encoding =
pd.get_dummies(df,columns=categorical_columns,drop_first=True)
encoding.head()
```

|   | Inches | Ram | Price       | ScreenWidth | ScreenHeight | Company_Apple |
|---|--------|-----|-------------|-------------|--------------|---------------|
| \ |        |     |             |             |              |               |
| 0 | 13.3   | 8   | 71378.6832  | 2560        | 2560         | True          |
| 1 | 13.3   | 8   | 47895.5232  | 1440        | 1440         | True          |
| 2 | 15.6   | 8   | 30636.0000  | 1920        | 1920         | False         |
| 3 | 15.4   | 16  | 135195.3360 | 2880        | 2880         | True          |
| 4 | 13.3   | 8   | 96095.8080  | 2560        | 2560         | True          |

|   | Company_Asus | Company_Chuiwi | Company_Dell | Company_Fujitsu | ... | \   |
|---|--------------|----------------|--------------|-----------------|-----|-----|
| 0 | False        | False          | False        | False           | ... | ... |
| 1 | False        | False          | False        | False           | ... | ... |
| 2 | False        | False          | False        | False           | ... | ... |
| 3 | False        | False          | False        | False           | ... | ... |
| 4 | False        | False          | False        | False           | ... | ... |

|   | Weight_4kg | Weight_5.4kg | Weight_5.8kg | Weight_6.2kg | Weight_7.2kg |
|---|------------|--------------|--------------|--------------|--------------|
| \ |            |              |              |              |              |
| 0 | False      | False        | False        | False        | False        |
| 1 | False      | False        | False        | False        | False        |
| 2 | False      | False        | False        | False        | False        |
| 3 | False      | False        | False        | False        | False        |
| 4 | False      | False        | False        | False        | False        |

|                  | Weight_8.23kg | Weight_8.4kg | Weight_? | CpuBrand_Intel |
|------------------|---------------|--------------|----------|----------------|
| CpuBrand_Samsung |               |              |          |                |
| 0                | False         | False        | False    | True           |
| False            |               |              |          |                |
| 1                | False         | False        | False    | True           |

|       |       |       |       |      |
|-------|-------|-------|-------|------|
| False |       |       |       |      |
| 2     | False | False | False | True |
| False |       |       |       |      |
| 3     | False | False | False | True |
| False |       |       |       |      |
| 4     | False | False | False | True |
| False |       |       |       |      |

[5 rows x 512 columns]

```

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

def evaluate_model(model, X_test, y_test, X_train, y_train):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    abs = mean_absolute_error(y_test, y_pred)
    sqr = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    print(f"Model: {model.__class__.__name__}")
    print(f" Mean Absolute Error: {abs: .2f}")
    print(f" Mean Square Error: {sqr: .2f}")
    print(f" R-Squared: {r2: .2f}")
    print('_' * 30)

    return abs, sqr, r2

```

```
lr = LinearRegression()
evaluate_model(lr, X_test, y_test, X_train, y_train)
```

Model: LinearRegression

Mean Absolute Error: 38418054892521.20

Mean Square Error: 29241147447119535695990882304.00

R-Squared: -26501020985635262464.00

---

(38418054892521.195, 2.9241147447119536e+28, -2.6501020985635262e+19)

Random Forest

```
rf = RandomForestRegressor(random_state=42)
evaluate_model(rf, X_test, y_test, X_train, y_train)
```

Model: RandomForestRegressor

Mean Absolute Error: 9679.64

Mean Square Error: 183796936.77

R-Squared: 0.83

---

(9679.636134245782, 183796936.7715833, 0.8334262878265054)

```
gb = GradientBoostingRegressor(random_state=42)
evaluate_model(gb, X_test, y_test, X_train, y_train)
```

Model: GradientBoostingRegressor

Mean Absolute Error: 11184.26

Mean Square Error: 226761900.16

R-Squared: 0.79

---

(11184.256581997686, 226761900.16170156, 0.7944874808420094)

## Hyperparameter Tuning

Hyperparameter Tuning for Random Forest

```
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5,
    scoring='neg_mean_squared_error', n_jobs=-1, verbose=2)
grid_search.fit(X_train, y_train)
print(f"Best Parameters: {grid_search.best_params_}")
```

```
print(f"Best CV Score: {-grid_search.best_score}")
best_model = grid_search.best_estimator_
evaluate_model(best_model, X_test, y_test, X_train, y_train)
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

## Hyperparameter Tuning for Gradient Boosting

```
param_grid_gb = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 4, 5],
    'subsample': [0.8, 0.9, 1.0]
}

# Setup the GridSearchCV
grid_search = GridSearchCV(estimator=gb, param_grid_gb=param_grid_gb,
cv=5, scoring='neg_mean_squared_error', n_jobs=-1, verbose=2)

# Fit the model
grid_search.fit(X_train, y_train)

# Print best parameters and best score
print(f"Best Parameters: {grid_search.best_params}")
print(f"Best CV Score: {-grid_search.best_score}")

# Evaluate the best model
best_model = grid_search.best_estimator_
evaluate_model(best_model, X_test, y_test, X_train, y_train)
```

## Real-time Predictions (Assuming Flask API):

```
from flask import Flask, request, jsonify
import joblib

app = Flask(__name__)
model = joblib.load('best_rf_model.pkl')

@app.route('/predict', methods=['POST'])
def predict():
    data = request.get_json()
    features = [data['RAM'], data['Storage'], data['Weight'],
data['Brand_Dell'], data['Brand_HP'], data['Brand_Lenovo']]
    features_scaled = scaler.transform([features])
    prediction = model.predict(features_scaled)[0]
    return jsonify({'predicted_price': prediction})

if __name__ == '__main__':
    app.run(debug=True)
```

```
## Interpretability and Insights
```

```
shap.initjs()
explainer = shap.Explainer(best_rf_model)
shap_values = explainer.shap_values(X_test_scaled)

shap.summary_plot(shap_values, X_test_scaled, plot_type='bar',
show=False)
plt.title('SHAP Values for Feature Importance')
plt.show()
```

```
-----
-----
ModuleNotFoundError                                Traceback (most recent call
last)
Cell In[1], line 1
----> 1 import shap
      3 shap.initjs()
      4 explainer = shap.Explainer(best_rf_model)
```

```
ModuleNotFoundError: No module named 'shap'
```

```
## Client Presentation (Using Plotly for Interactive Dashboards)
```

```
import plotly.express as px
```

```
# Create interactive dashboards
```

```
fig = px.scatter(df, x='RAM', y='Price', color='Brand',
trendline='ols', title='RAM vs. Price by Brand')
fig.show()
```